



Making C++ Run Better with

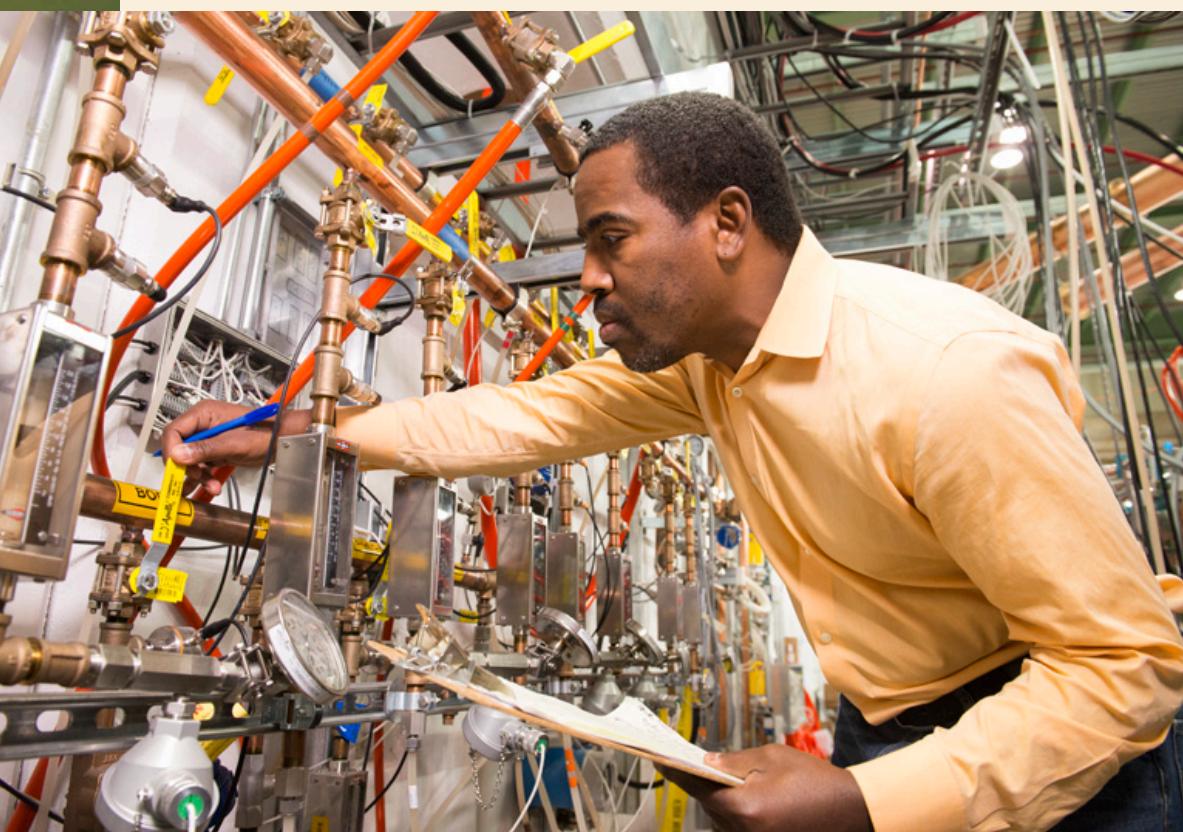
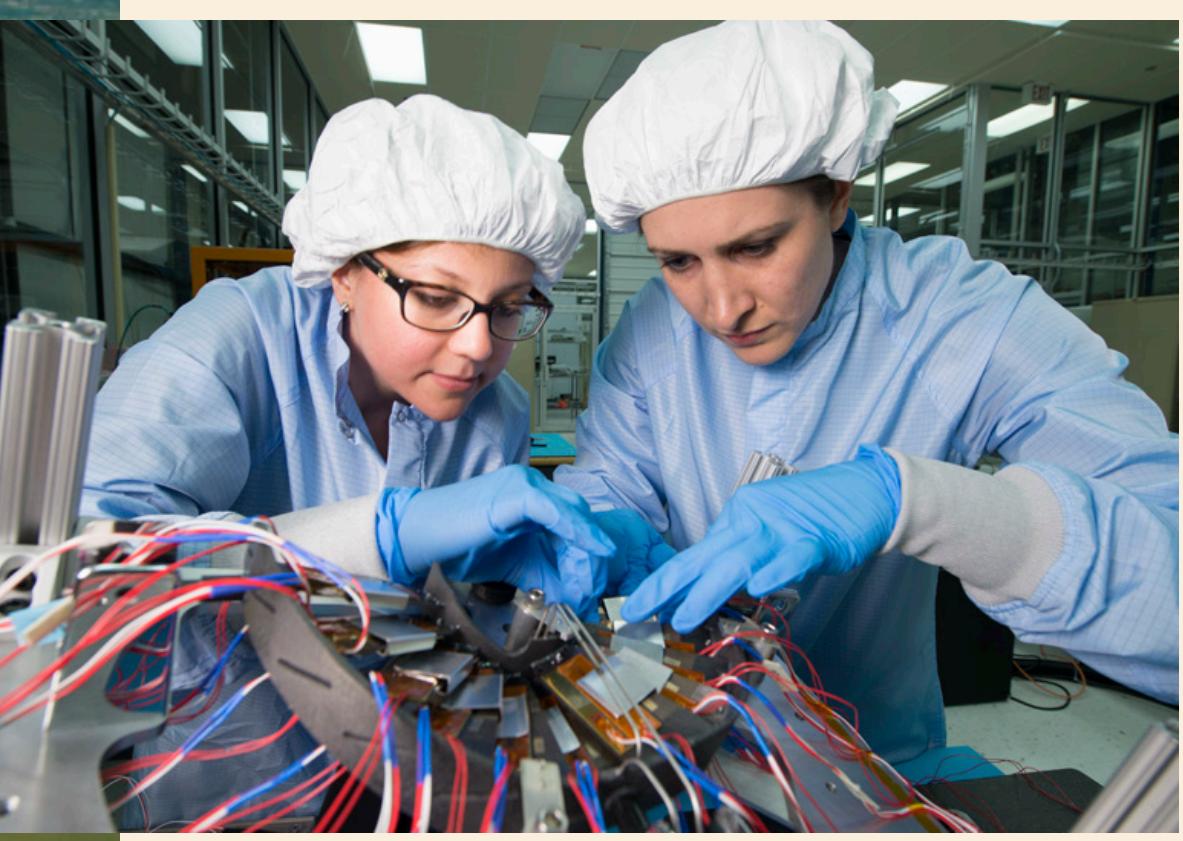
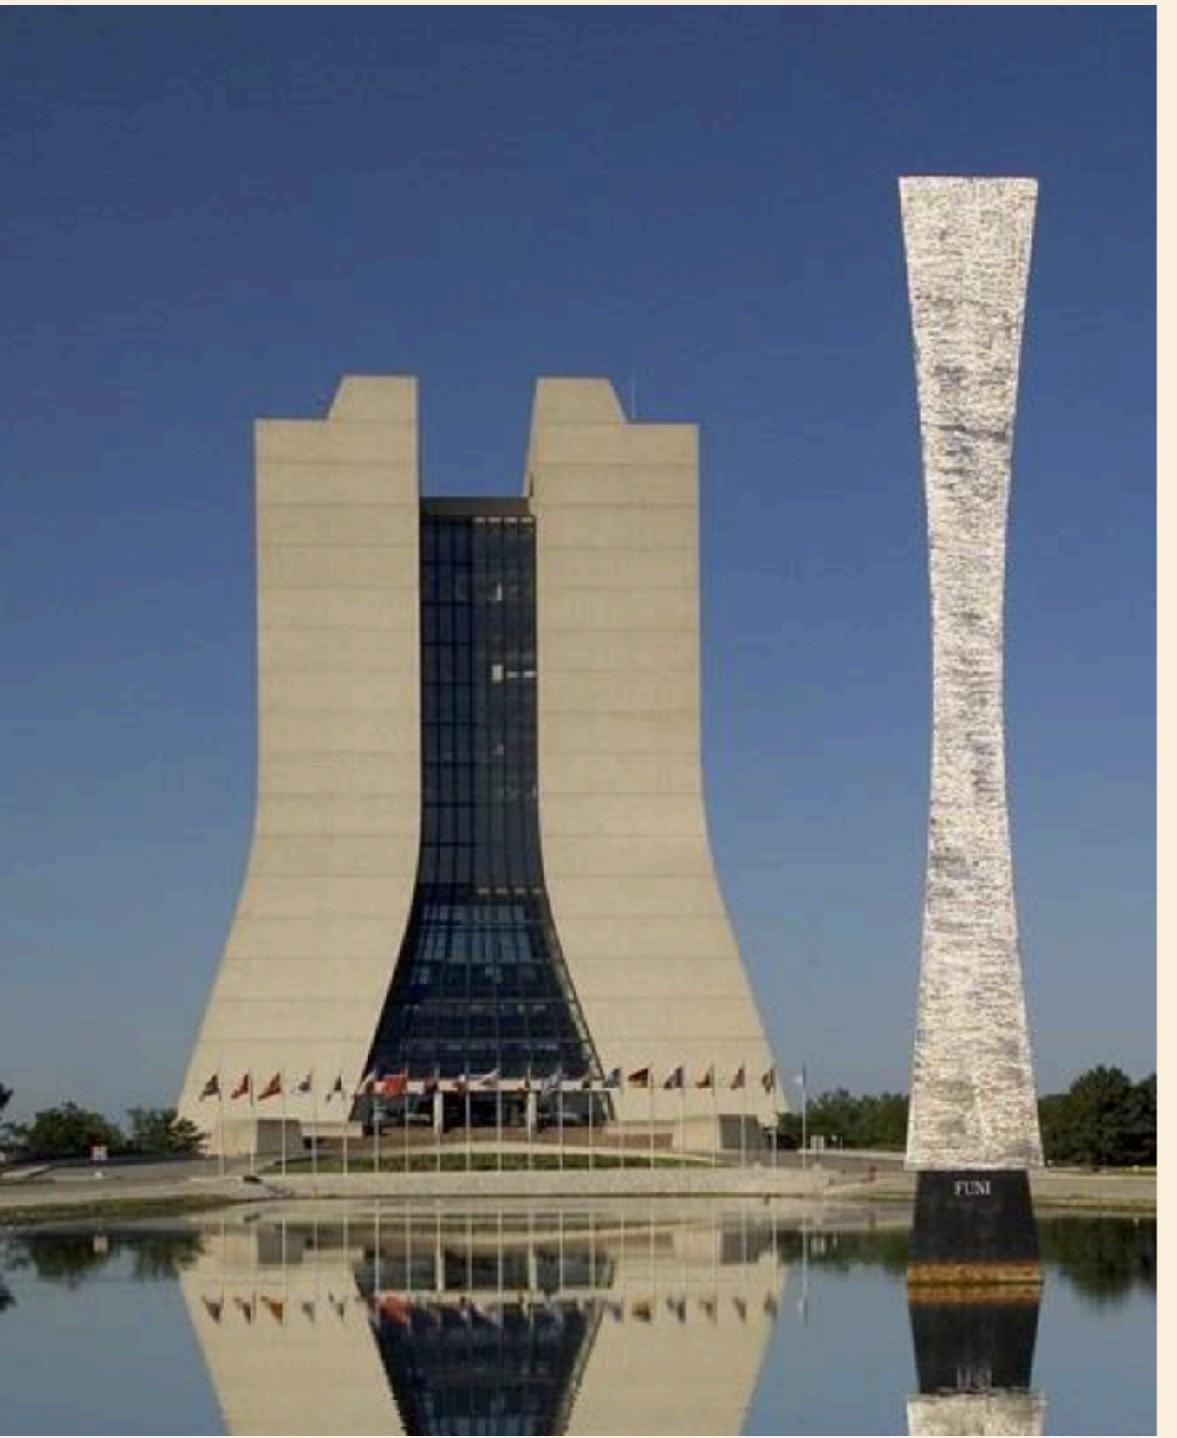
Adam Lyon (Fermilab Scientific Computing Division)

SATR**DAYS Chicago**

27 April 2019

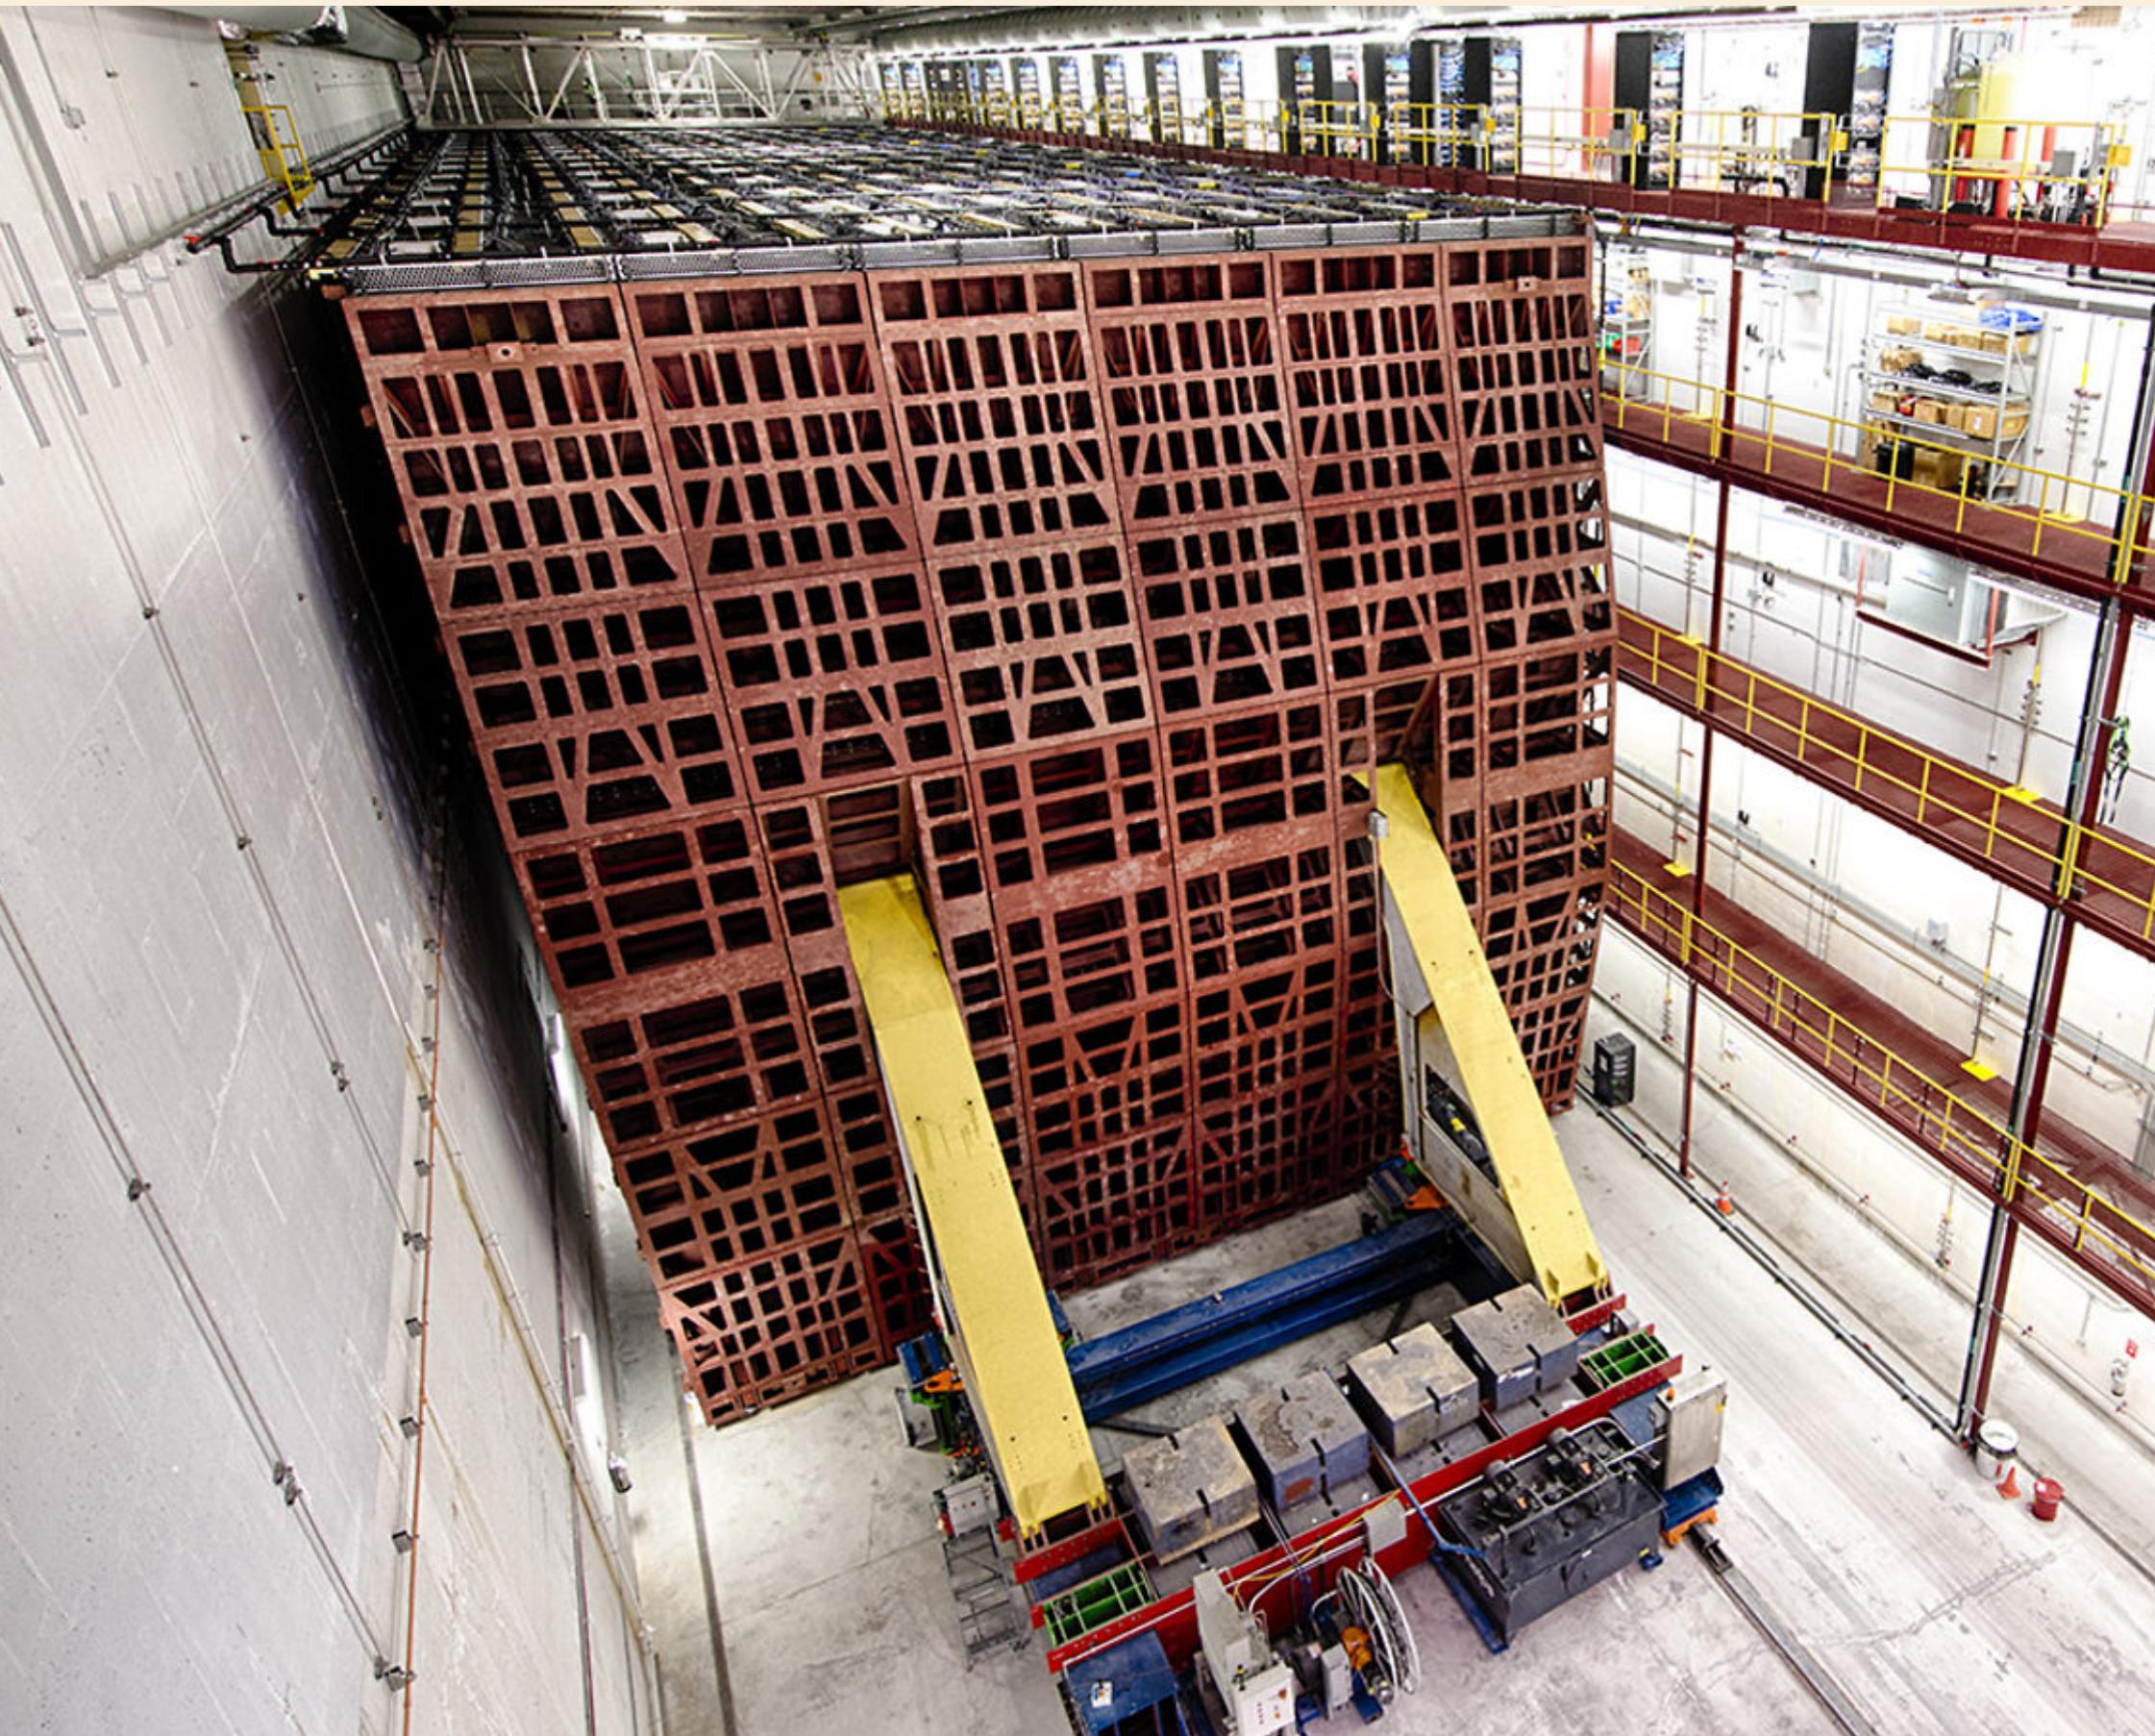
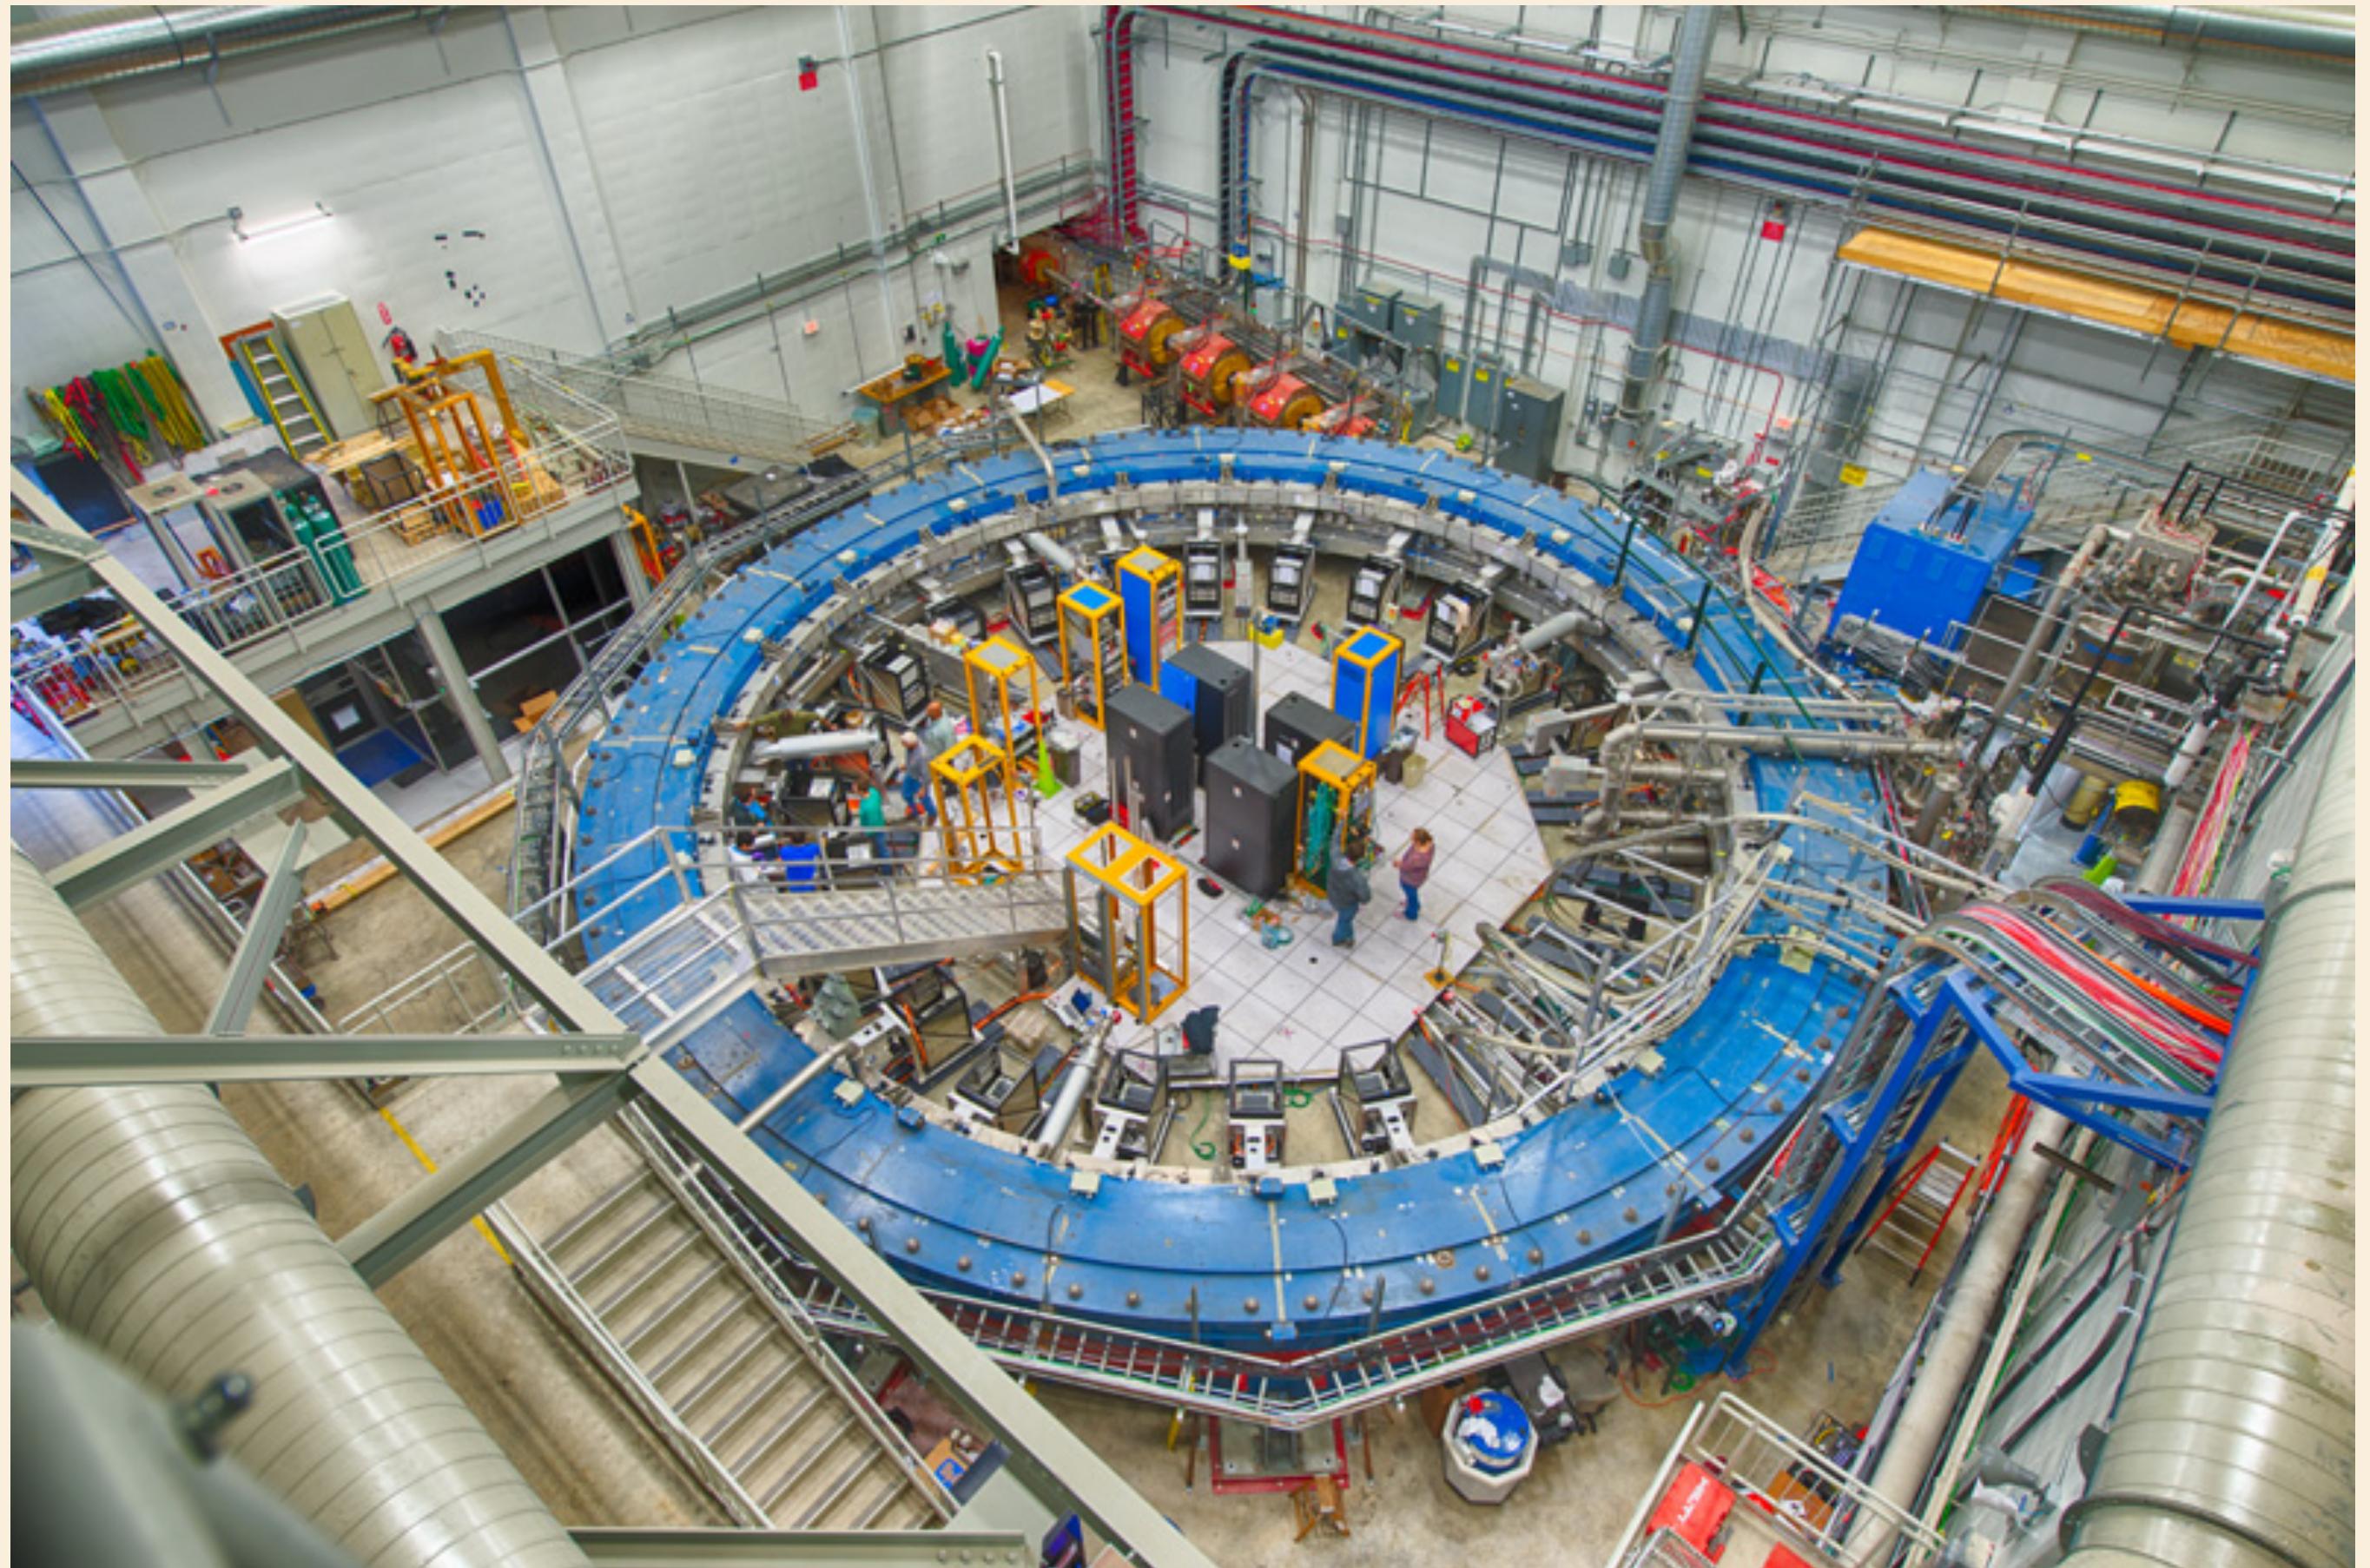
Fermi National Accelerator Laboratory

Fermilab is America's Particle Physics and Accelerator Laboratory



We build huge machines to understand how our universe works

We do science! We use machines to search for smallest things we can find in nature and learn more about them (subatomic particles)





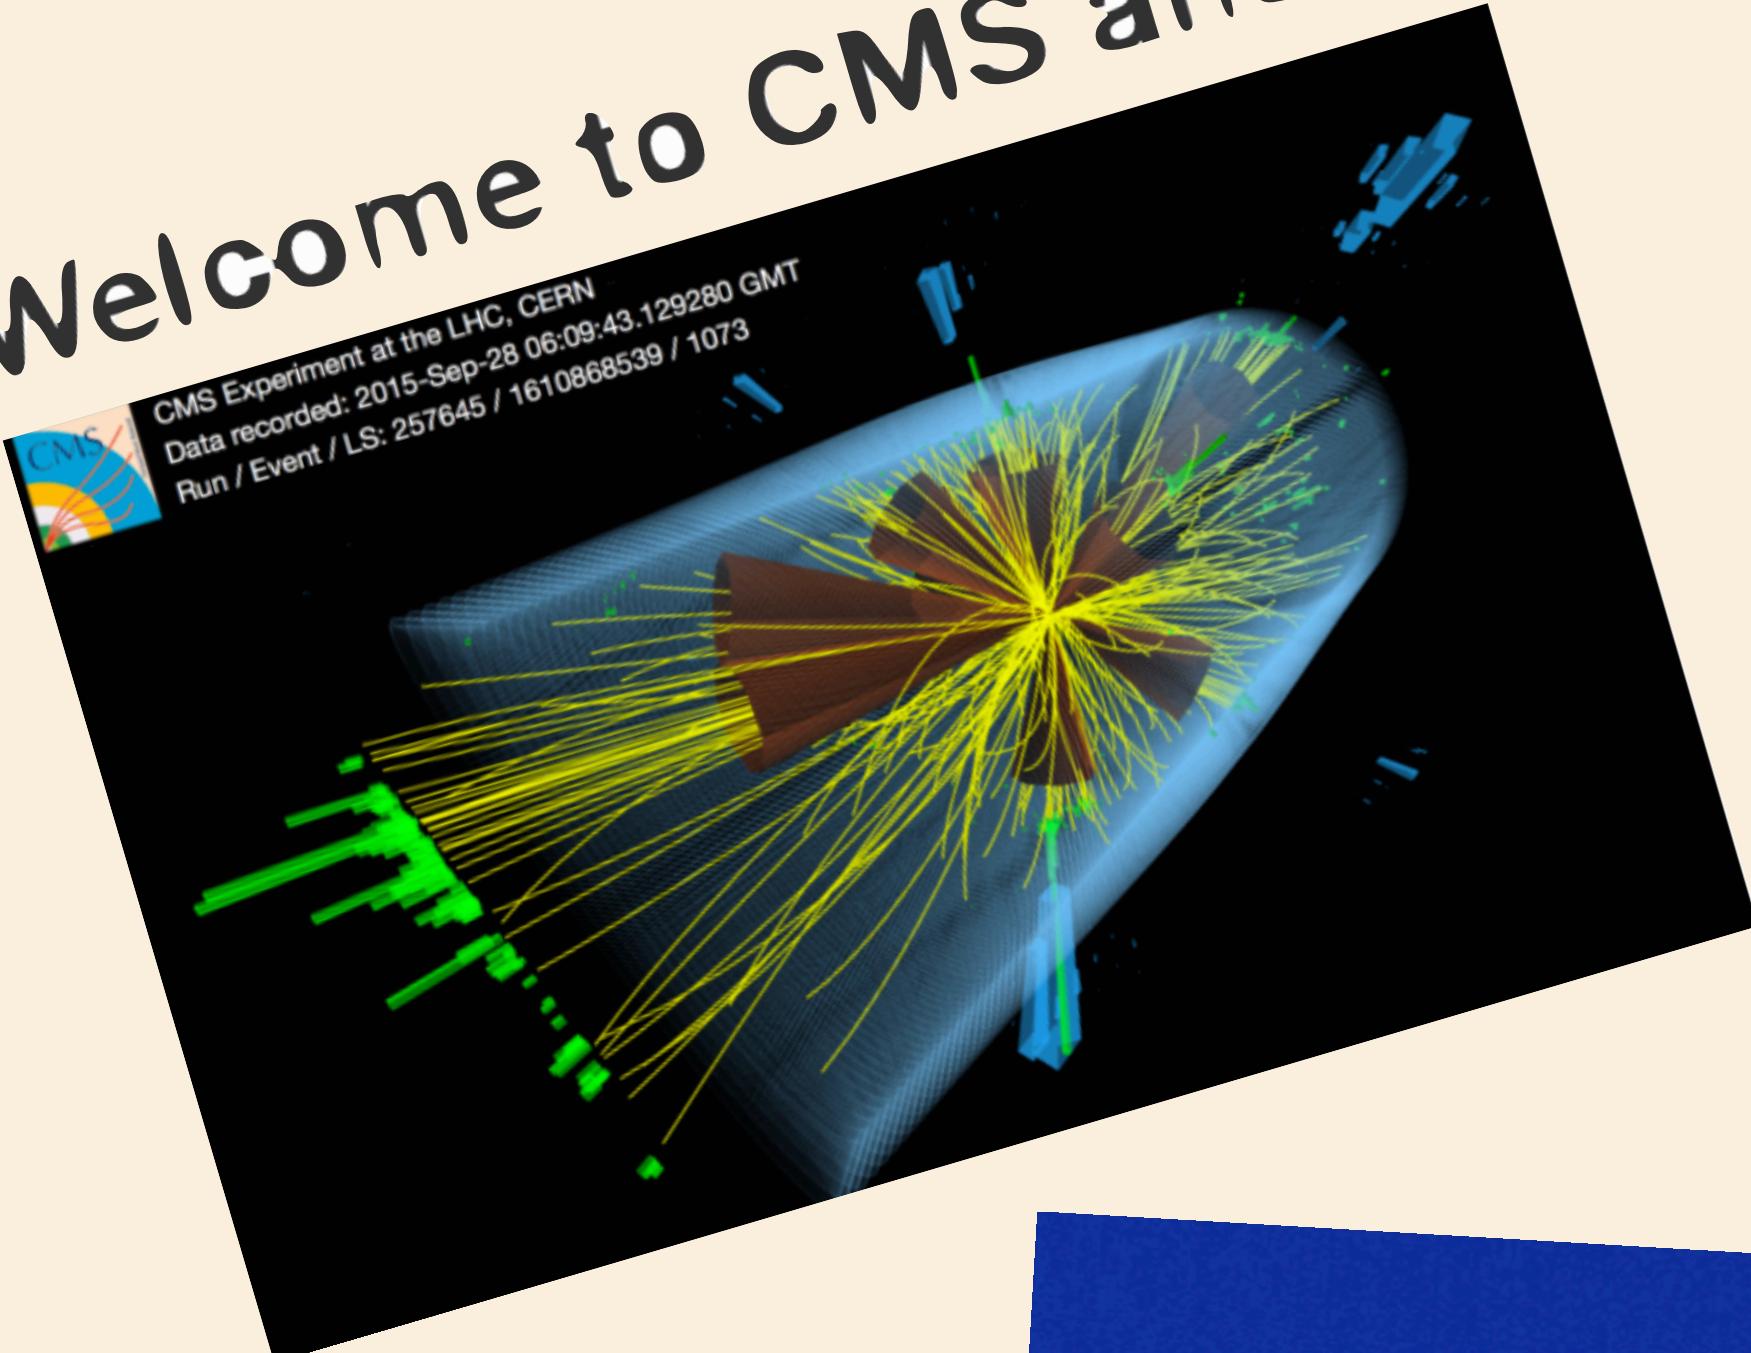
**And we write lots
of code!**

**Mostly written by many
physicists who are *not*
software experts**

**Except for difficult
infrastructure code
written by a few
computing professionals
who are software experts**

Lots and lots and lots of code...

Welcome to CMS and CMSSW



Millions of lines of
C++

Process raw data from
experiments into
physics objects
(reconstruction)

Simulations

Analyze data

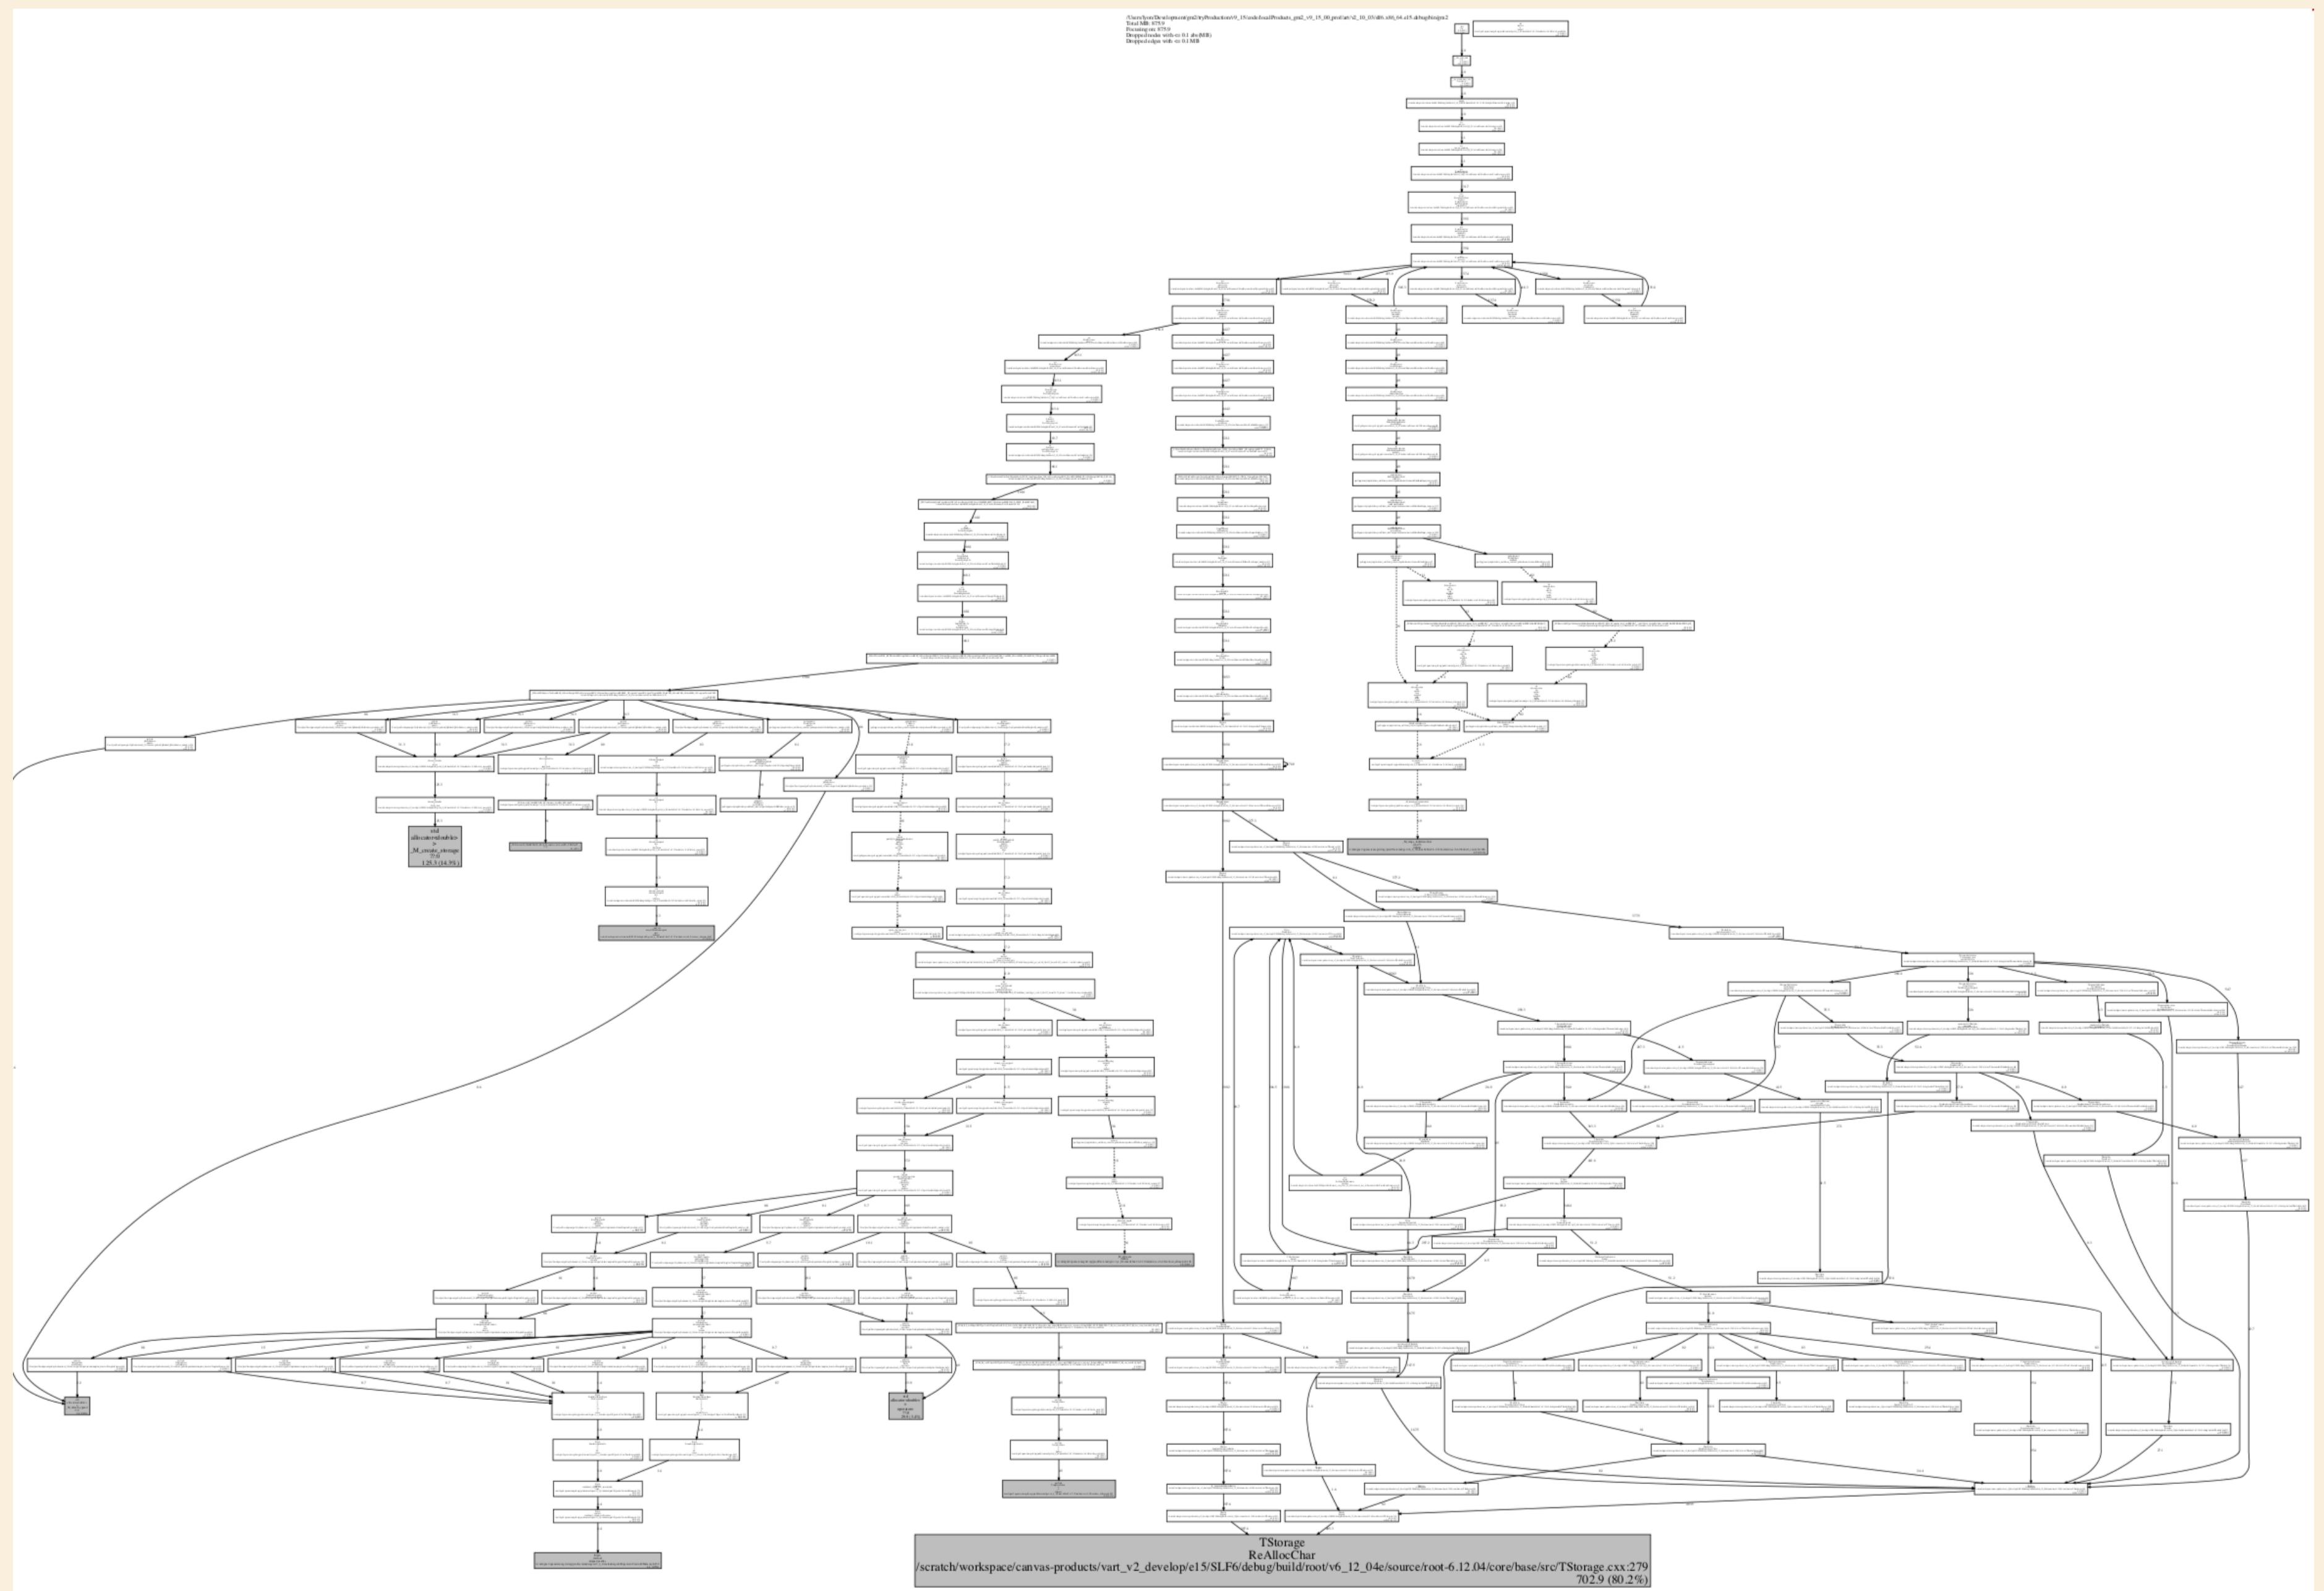
All open source

Things can go wrong

Code takes too long to run

Code uses too much memory

Run a profiler, like Google's gperftools, and get this...



Now what?



Zooming in helps, a little

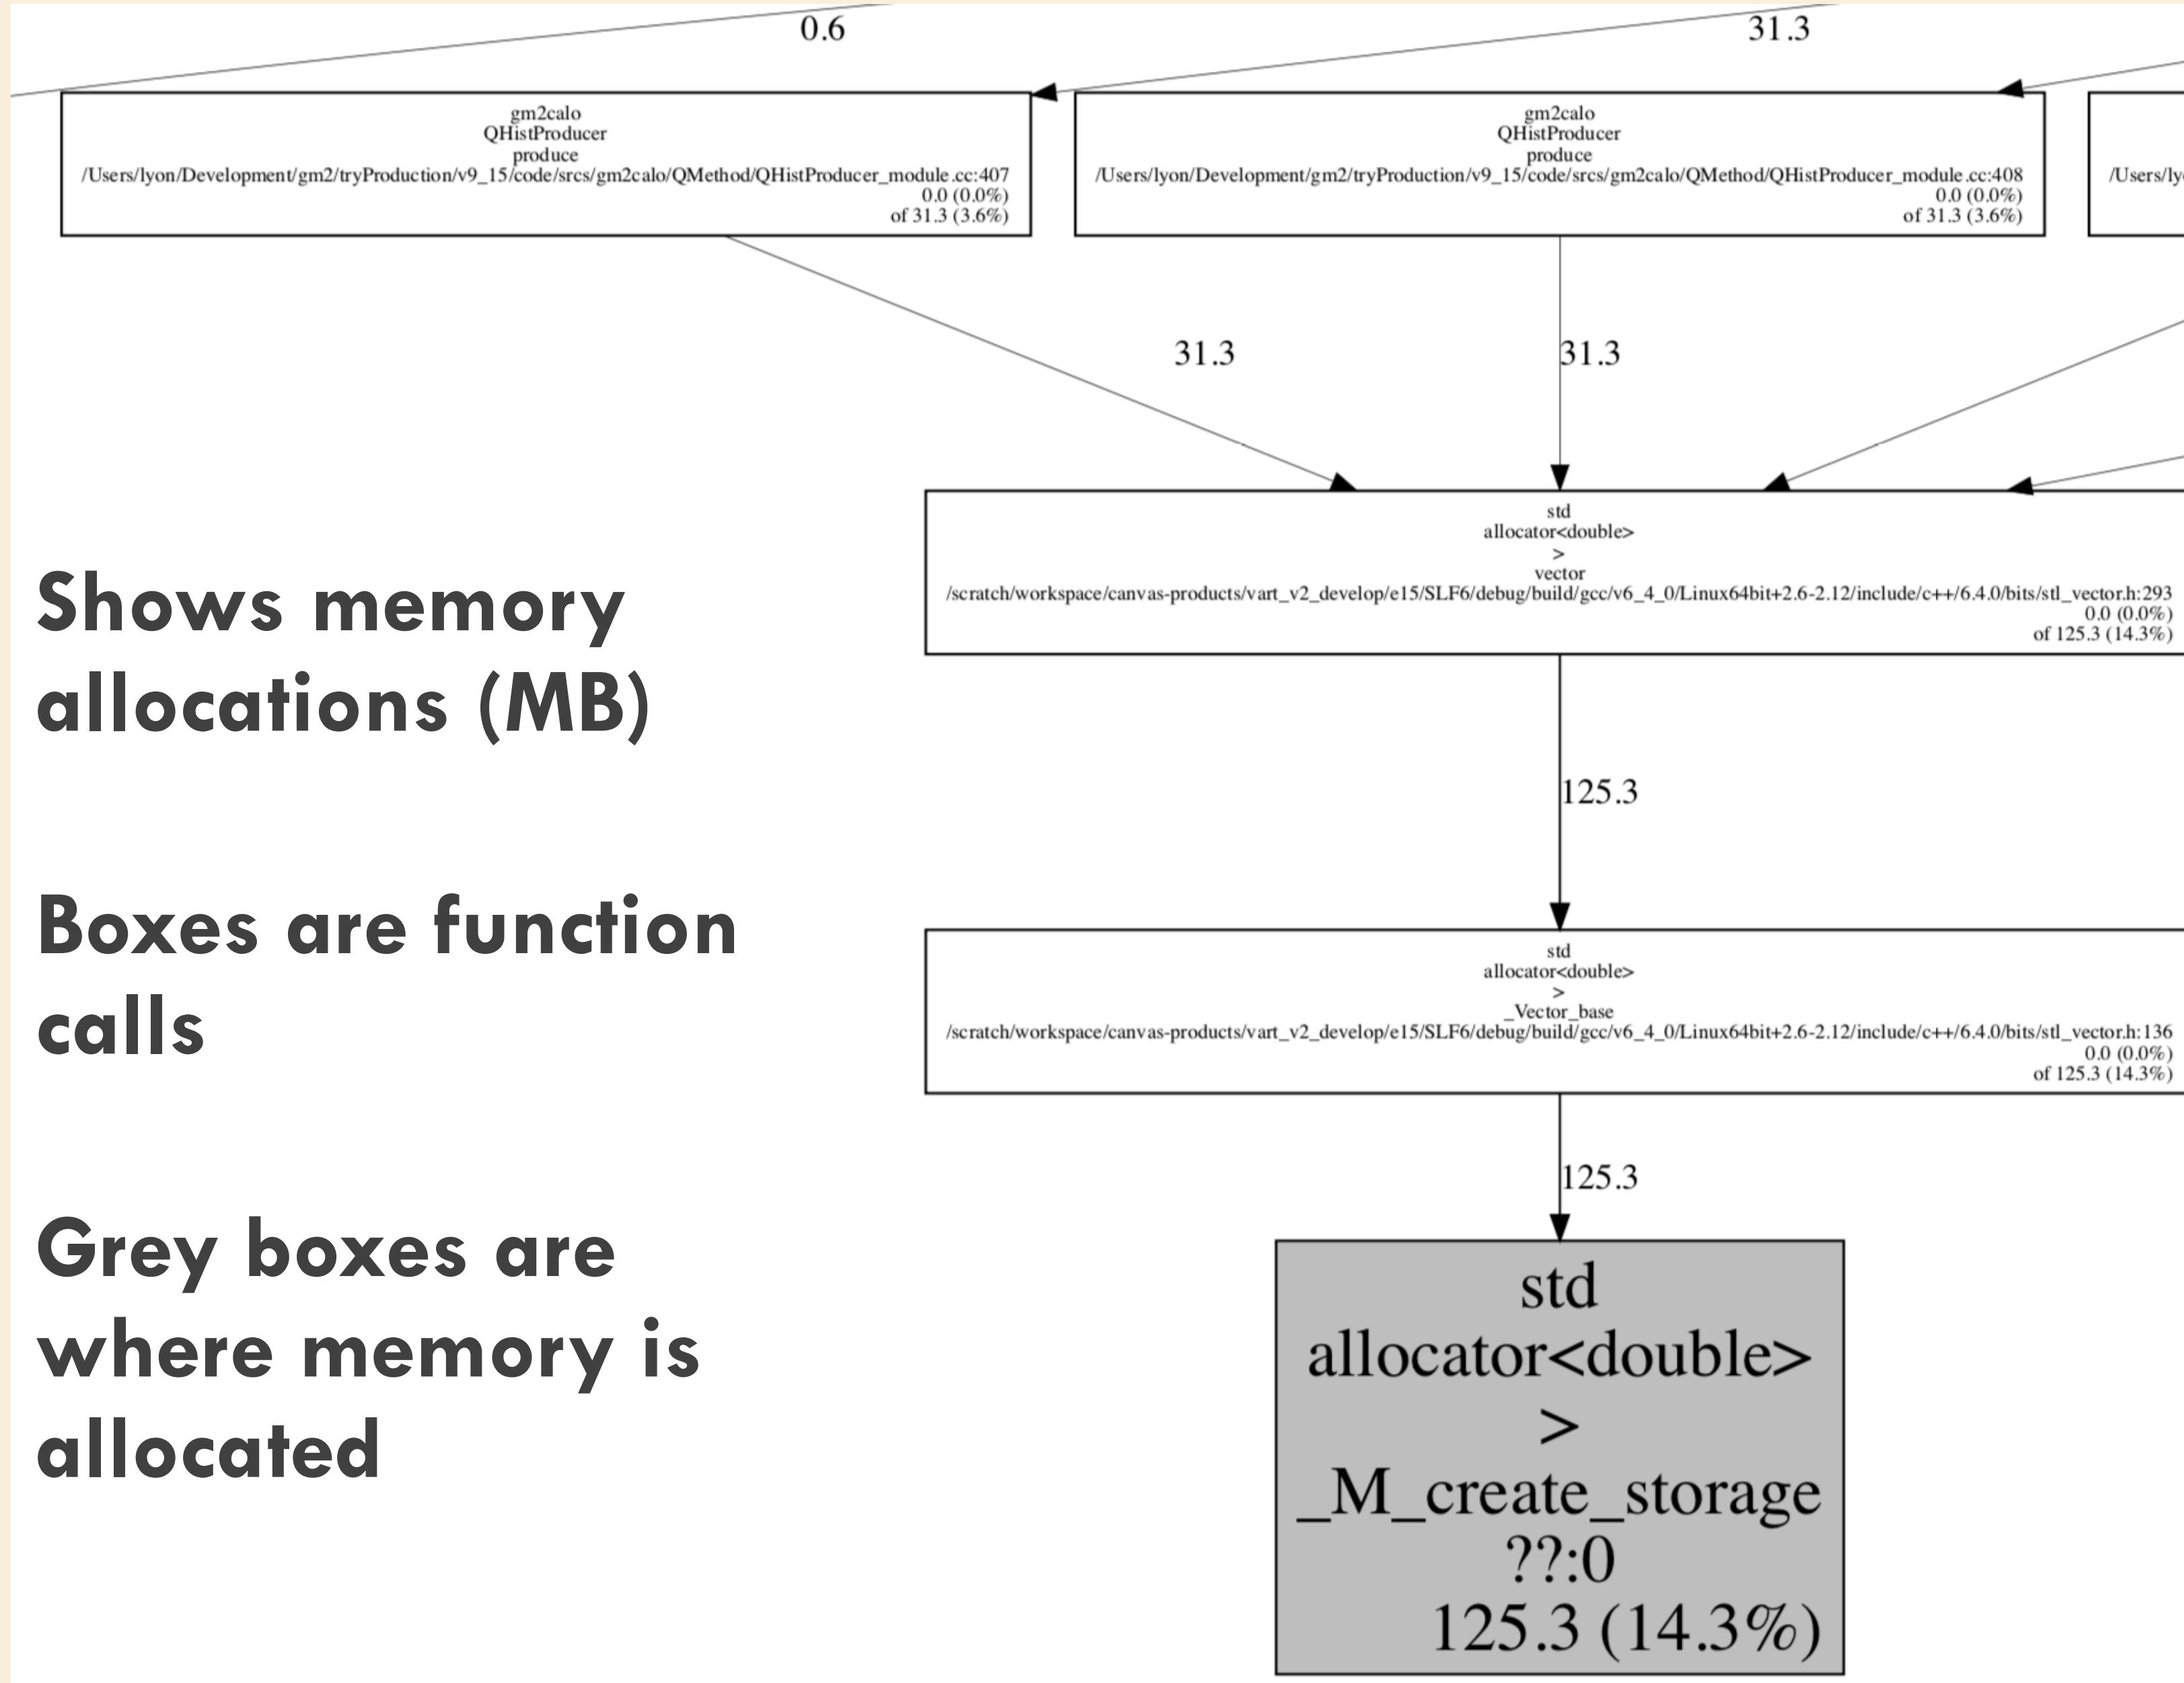
How to analyze programmatically?

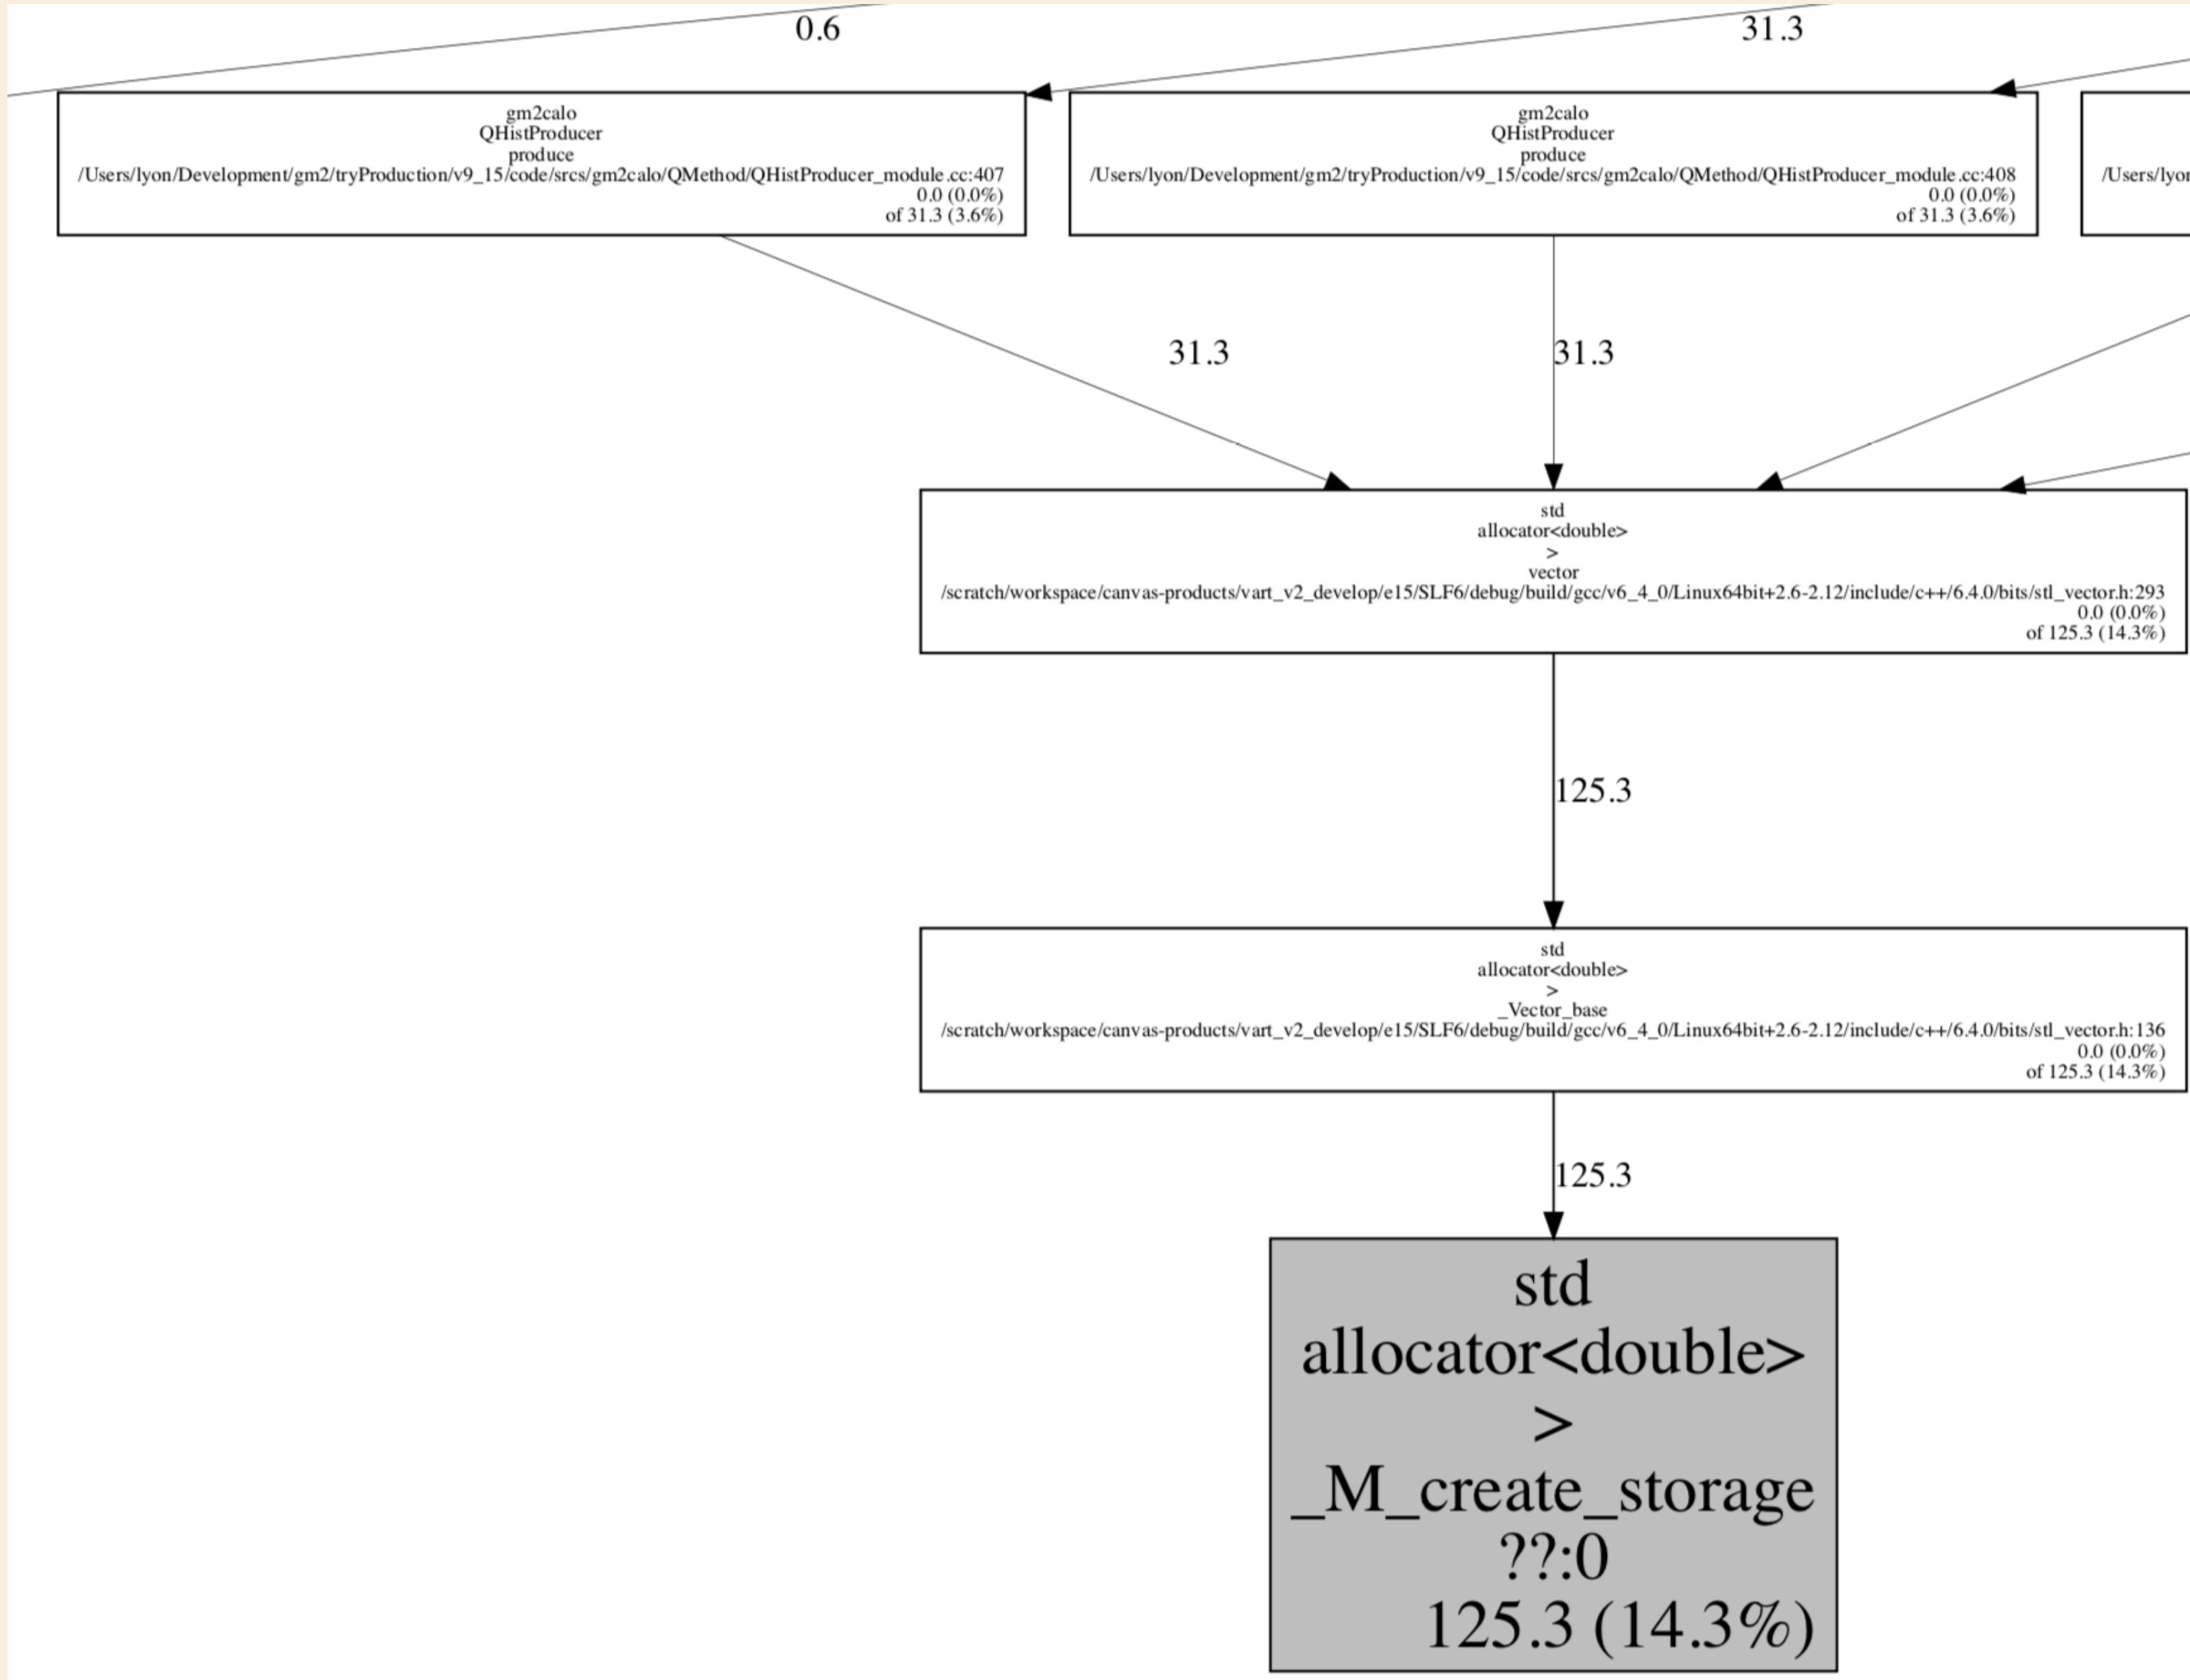
How to analyze reproducibly?

**Shows memory
allocations (MB)**

**Boxes are function
calls**

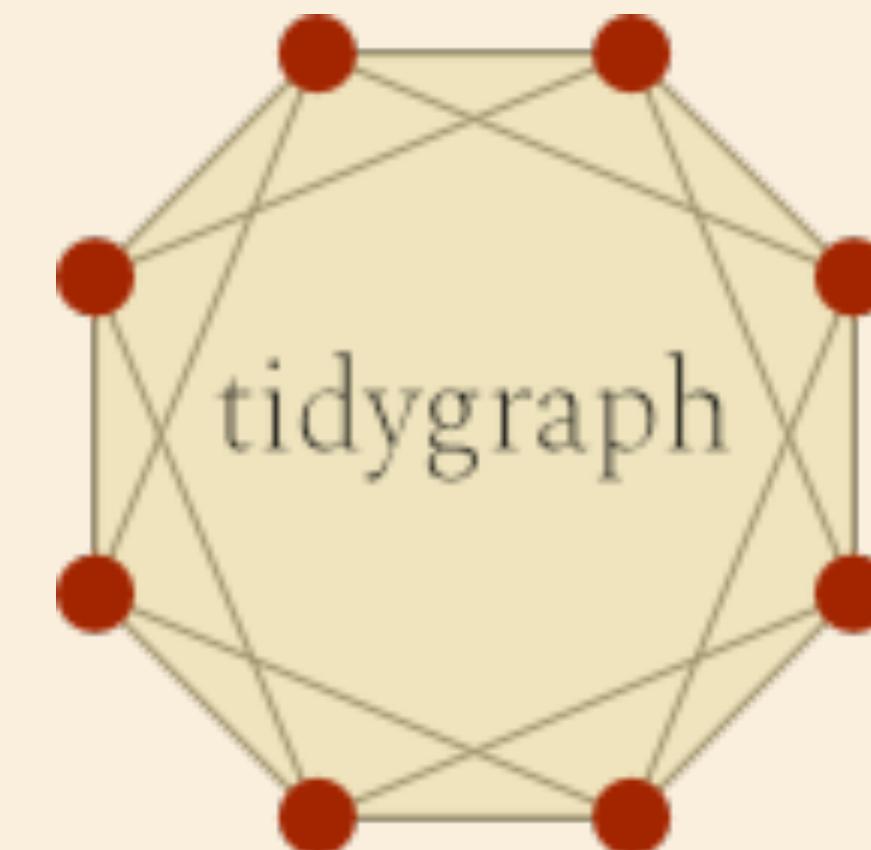
**Grey boxes are
where memory is
allocated**





This looks like a
directed graph

R can do that!



Analyzing gperftools output with Tidygraph

gperftools output → Tidygraph

1. `$ dot -Tdot_json out.dot > out.dot.json # Convert dot file into json`
2. `jsonlite::fromJSON(out.dot.json) # Read into R`
3. *Extract nodes (boxes) and edges (lines) to data frames and “fix them up”*
4. `g ← tidygraph::tbl_graph(nodes=nodes, edges=edges) # Make Tidygraph`

Do stuff: Find the memory allocators (bottom of the graphs)

```
g %>%
  filter(node_is_sink()) %>%
  as_tibble() %>%
  select(index, name, totalMem, totalPerc) %>%
  arrange(-totalMem)
```

index	name	totalMem	totalPerc
<int>	<chr>	<dbl>	<dbl>
1	TStorage ReAllocChar	702.9	80.2
71	std allocator<double> > _M_create_storage	125.3	14.3
109	std allocator<double> > operator=	29.9	3.4
148	_M_range_initialize<char const*> (inline)	4.9	0.6
152	Eigen internal aligned_malloc	4.4	0.5
155	_M_allocate (inline)	3.8	0.4
169	std allocator<double> > _M_default_append	3.4	0.4
200	gm2calo TemplateFitter> > > allocate	0.5	0.1
202	TBufferFile WriteFastArray@176fe6	0.5	0.1
205	TBufferFile WriteFastArray@176c32	0.3	0.0

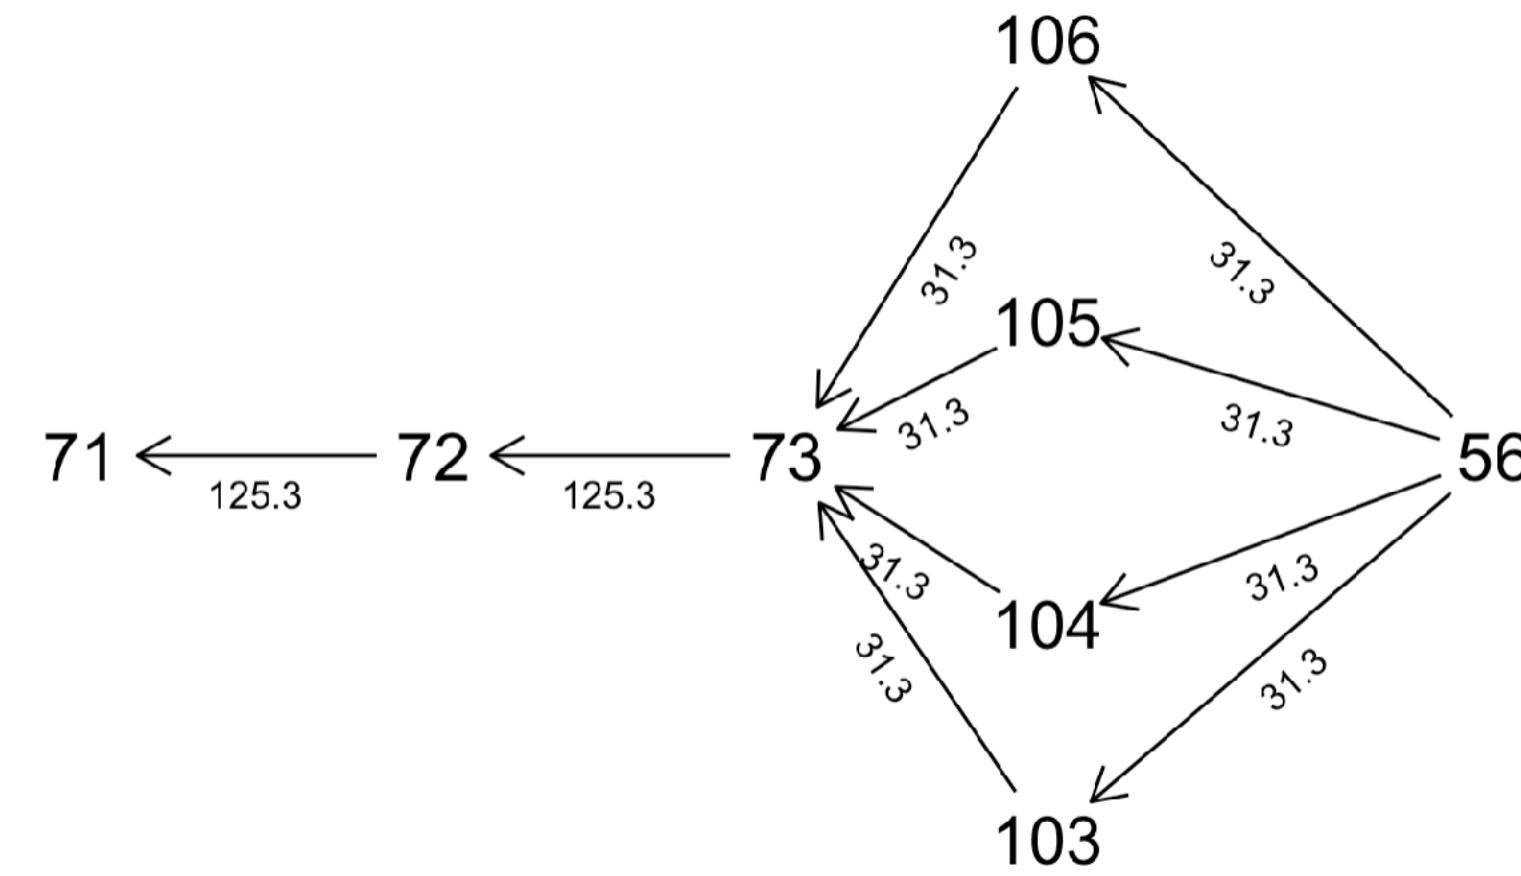
Traverse up from one of them

```
g %>%
  mutate(dist = dfs_dist(71, mode="in")) %>%
  filter(!is.na(dist)) %>%
  as_tibble() %>%
  arrange(dist) %>%
  select(index, name, file, totalMem, dist)
```

This is **readable**,
programmable,
reproducible

index	name	file	totalMem	dist
71	std::allocator<double>::_M_create_storage	???:0	125.3	0
72	std::allocator<double>::_Vector_base	stl_vector.h:136	125.3	1
73	std::allocator<double>::vector	stl_vector.h:293	125.3	2
103	gm2calo::QHistProducer::produce	QHistProducer_module.cc:406	31.3	3
104	gm2calo::QHistProducer::produce	QHistProducer_module.cc:407	31.3	3
105	gm2calo::QHistProducer::produce	QHistProducer_module.cc:408	31.3	3
106	gm2calo::QHistProducer::produce	QHistProducer_module.cc:409	31.3	3
56	_ZN3art10EDProducer7doEventERNS_14EventPrincipal...	EDProducer.cc:25	168.1	4

Traverse up from one of them



Make easier to read
pictures too

index	name	file	totalMem	dist
71	std::allocator<double>::_M_create_storage	???:0	125.3	0
72	std::allocator<double>::_Vector_base	stl_vector.h:136	125.3	1
73	std::allocator<double>::vector	stl_vector.h:293	125.3	2
103	gm2calo::QHistProducer::produce	QHistProducer_module.cc:406	31.3	3
104	gm2calo::QHistProducer::produce	QHistProducer_module.cc:407	31.3	3
105	gm2calo::QHistProducer::produce	QHistProducer_module.cc:408	31.3	3
106	gm2calo::QHistProducer::produce	QHistProducer_module.cc:409	31.3	3
56	_ZN3art10EDProducer7doEventERNS_14EventPrincipal...	EDProducer.cc:25	168.1	4

Here are the culprits (allocates ~15% of total memory usage)

```
405     xtalN_[iCalo].push_back(iChan);  
406     e_Traces_[iCalo].push_back(std::vector<double>(nBins_[iCalo], 0));  
407     eSqSum_Traces_[iCalo].push_back(std::vector<double>(nBins_[iCalo], 0));  
408     rawSum_Traces_[iCalo].push_back(std::vector<double>(nBins_[iCalo], 0));  
409     pedSum_Traces_[iCalo].push_back(std::vector<double>(nBins_[iCalo], 0));
```

Takeaways:

Make interactive tasks programmable and reproducible

Easy to document. Easy to share. Easy for other people to try and extend.

See <http://bit.ly/gperftoolsR>

Fermilab is more than physics — open to the public — come visit!

